



Optimal Parallel Randomized Algorithms for Sparse Addition and Identification

Advanced algorithms and parallel programming
Valentina Ionata Pietro Ferretti

General problem

➤ *Parallel addition of n numbers in shared memory models*

Standard algorithms are not sensitive to properties of the input, thus randomized algorithms can beat standard parallel algorithms in specific cases:

e.g. let's consider the case when there is **a large amount of zero operands** among the numbers to be added

Proposed Solution

Randomized parallel algorithm:

- the number of non-zero operands must be much smaller than the number of zeros among the numbers to be added
- we suppose we are using a CRCW PRAM with unlimited memory and unlimited number of processors
- time complexity $O(\log m)$, space complexity $O(m)$ w.h.p.

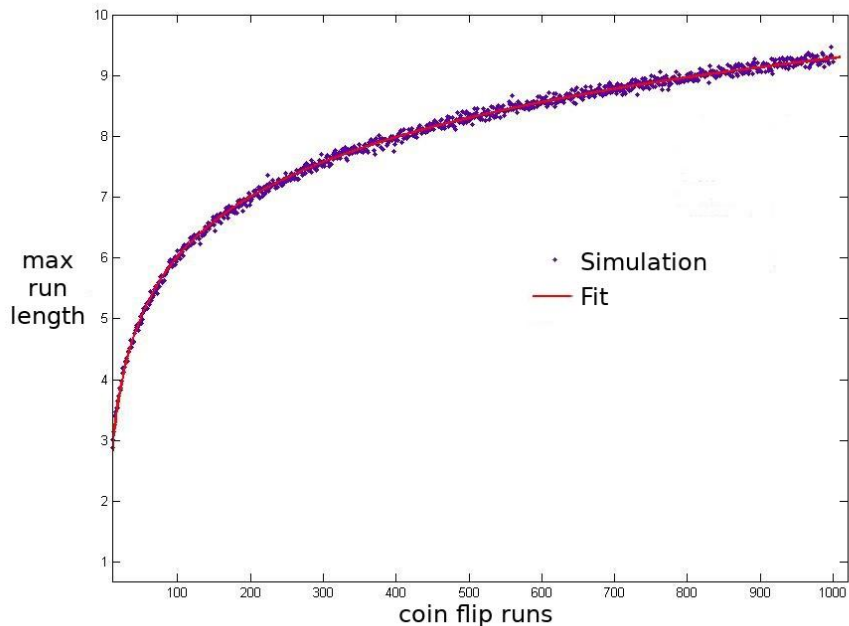
The Algorithm

It consists of two phases:

- **Estimation** of the number m of non-zero terms
- **Parallel addition** on a smaller set of inputs

Estimation - *First phase*

1. Each processor P_i has a number to add.
2. **PRODUCE-AN-ESTIMATE:**
every processor with $x_i \neq 0$ flips a coin until a head is found. A shared memory location will contain the *highest number of coin flips reached*.



Estimation - *First phase*

3. PRODUCE-AN-ESTIMATE is executed ***k*** times.
4. Then each P_i computes an estimation m' for the number of non-zero operands m using the following formula:

$$(1) \quad E \leftarrow (\log 2) \frac{E_1 + \dots + E_k}{k}$$

$$(2) \quad m' \leftarrow \exp(2) \cdot \exp(E) + d$$

where $d \geq 1$ is a constant.

Code (Estimation)

INITIALIZATION

Processor P_1 initializes a special shared memory location (CLOCK) to zero.
Then, each P_i executes $TIME_i \leftarrow 0$.

ESTIMATE

```
Processor  $P_i$ 
  if  $x \neq 0$ 
    then begin (1) Flip a fair coin (two-sided)
              (2 ) If the outcome is 'tail' then
                    begin (2a)  $TIME_i = TIME_i + 1$ 
                          (2b)  $CLOCK \leftarrow TIME_i$ ;
                          (2c) goto(1)
                    end
              end
    end
```


#all processors wait for everyone to be finished

Each P_i with $x_i \neq 0$ reads CLOCK and makes its value to be the current estimate.

Addition - *Second phase*

1. Initialization:

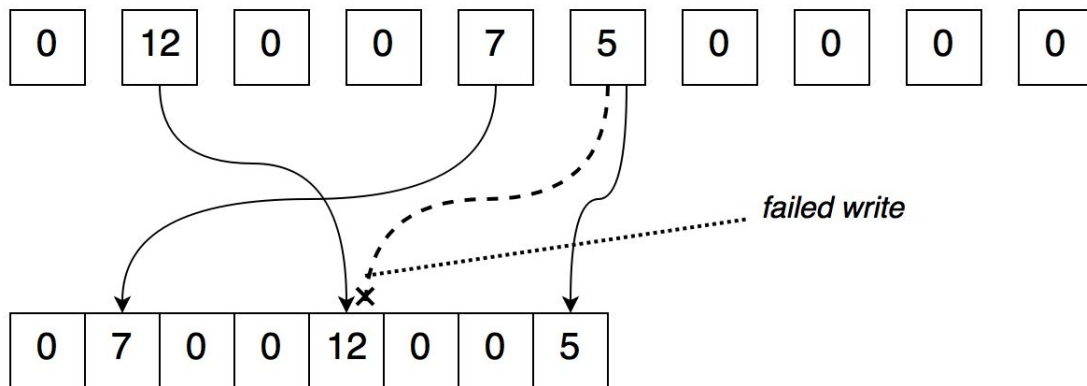
$a \cdot m'$ memory locations are set to zero. ($a \geq 4$ constant)



Addition - *Second phase*

2. Memory Marking:

each processor with a non-zero number tries to write it to a random empty location in the zeroed space, until they all succeed.

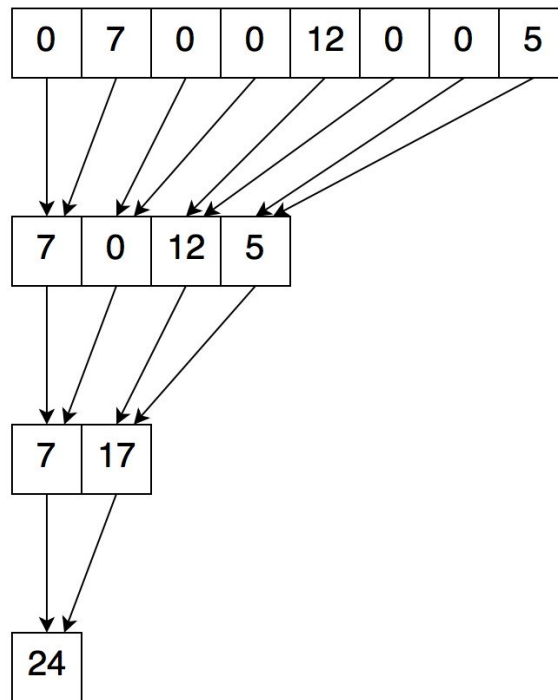


Addition - *Second phase*

3. **Addition:**

standard parallel addition of all the $a \cdot m'$ numbers (zero and non-zero).

The result is of course the sum of all the non-zero numbers.



Code (Addition)

INITIALIZATION

In one parallel step, processors initialize a $\cdot m' + 2$ shared memory locations to zero, by executing: “Processor P_i writes a zero to $M(j)$, if $j < am' + 2$ ”. Then, they all execute $TIME_j \leftarrow 0$.

MEMORY MARKING

Processor P_j

IF $x_j \neq 0$ then

BEGIN (1) Select y equiprobably at random from $\{ 1, 2, \dots, am' \}$

(2) $TIME_j \leftarrow TIME_j + 1$

(3) Read $M(y)$; $TIME_j \leftarrow TIME_j + 1$

(4) If $M(y) = 0$ then write x_j into $M(y)$. Also, $TIME_j \leftarrow TIME_j + 1$

(5) If the “write” failed then

BEGIN (5a) write $TIME_j$ into $M(am' + 1)$

(5b) go to (1)

END

END

Code (Addition)

The following part is executed by P_j with $x_j = 0$ and by “successful” P_j with $x_j \neq 0$.

(6) Read $M(am' + 1)$ into a local variable $R1$

(7) Wait for 8 steps

(8) Read $M(am' + 1)$ into a local variable $R2$

(9) If $R1 \neq R2$, then go to (6)

ADDITION

(Processor P_j is assigned to location $M(j)$, $1 \leq j \leq am'$) From this point on, processors P_j (where $1 \leq j \leq am'$) perform a standard parallel addition of the numbers $M(1), \dots, M(am')$. In the i -th parallel step of the addition, processor P_j adds $M(j)$ and $M(j + 2^i)$ into $M(j)$, for $j = k \cdot 2^i + 1$, $k = 0, 1, 2, \dots, am'/2^i$.

Properties

- The sum is always correct
- The time complexity is in $O(\log m)$ and the space complexity in $O(m)$ (where m is the number of non-zero entries among the terms we need to sum) with high probability
- We can control the precision of the estimate, and thus the variance of the running time, by changing the amount of estimations of m to run (represented by the constant k)

Possible Applications

Marking algorithm:

- solves the processor identification problem for a strong WRAM in $O(m)$ parallel time with arbitrarily high probability.

Processor identification problem:

N processors are given, each keeping either a 0 or a 1. The problem is for each processor to find out which are the processors with the 1's.

Potential Problems

- The algorithm assumes a CRCW PRAM
 - we think this can be implemented using atomic compare-and-swap primitives
- We have few cores at our disposal (up to 4). It's not a given that the performance will be as good as the one predicted by the paper for n processors.
- We will have to find working values for k , a , d since they're not specified in the paper
- Since we will be summing fixed-width numbers we will have to consider overflows

Testing and Performance

Datasets for testing are easy to generate (random numbers).

Tests to perform:

- different amounts of numbers to sum
- different ratios of zero and non-zero terms
- different values of k , a , d

Comparison with other algorithms:

- performance gain w.r.t. a serial algorithm
- performance gain w.r.t. the standard parallel sum (parallel “lower bound”)

Planned Schedule

- Implementation:
 - serial algorithm
 - parallel “lower bound”
 - parallel randomized algorithm
- Testing
- Presentation